

ADAPT: A Self-Calibrating Proactive Autoscaler for Container Orchestration

Himanshu Singh Baghel

Department of Computer Engineering
[J.C. Bose University of Science and Technology]
Faridabad, India
baghelhimanshu2004@gmail.com

Abstract—Proactive autoscaling for containerized workloads depends on knowing the provisioning delay, i.e., the time between a scaling decision and the moment new capacity is ready to serve traffic. In practice, this cold-start duration can vary substantially across environments and even across consecutive scale-out events. We present ADAPT (Adaptive Duration Approximation for Predictive Timing), an online EWMA estimator that tracks cold-start duration at runtime. ADAPT feeds a dynamic planning horizon, FH-OPT, into a Model Predictive Controller (MPC) that optimizes replica counts over a rolling window. Together, these components form a closed-loop proactive autoscaling design that adapts its lookahead based on measured provisioning delay. Evaluated across three policies (MPC+LSTM, MPC+Prophet, HPA) and six workload archetypes with five random seeds, MPC+LSTM achieves below 5% SLA violation on all workloads, compared with 7–19% for reactive HPA and up to 28.7% for MPC+Prophet on bimodal traffic.

Index Terms—Autoscaling, Cold Start, Model Predictive Control, EWMA, Proactive Scaling, Kubernetes, SLA, Forecasting

I. INTRODUCTION

Modern containerized applications are routinely deployed on orchestration platforms like Kubernetes, where horizontal scaling directly determines both service reliability and infrastructure cost. The de facto standard for this task, the Kubernetes Horizontal Pod Autoscaler (HPA) [1], operates on a simple reactive principle: observe current resource utilization, compare it against a target threshold, and adjust the replica count accordingly. Every scaling decision is made in isolation, with no memory of prior load behavior and no foresight into where demand is heading.

This reactive paradigm has a well-known limitation in containerized environments: the startup gap. When a scaling decision is triggered, newly requested replicas do not become available instantly. Each container must boot, initialize its runtime, load application dependencies, and begin accepting traffic, a process that routinely takes between 30 and 300 seconds depending on the application and cloud provider [2], [3]. By the time HPA detects an overload condition and new replicas are ready, the system has already been degraded for the duration of that cold-start window.

The natural response is proactive autoscaling: predict future demand and initiate scaling decisions early enough that new replicas are ready before the load arrives. This idea is well established in the literature [4]–[8]. A 2025 survey of 47

autoscaling systems [9] found that proactive approaches consistently outperform reactive baselines on workloads with predictable demand patterns. However, the same survey identified a recurring limitation: most proactive systems, both academic and production, treat the container cold-start duration as a static configuration constant hardcoded at deployment time.

This assumption is problematic in practice. Cloud platform boot times are not fixed. AWS EC2 instance initialization varies by ± 40 –60% under real-world conditions; GKE pod cold-starts fluctuate by ± 20 –50% depending on node pressure, image layer caching, and cluster autoscaler behavior [2]. A proactive scaler that plans its horizon around a static $\Delta_{\text{cold}} = 120$ s will systematically under-provision when actual boot times reach 180 s, and waste resources when those same boots complete in 80 s. The planning horizon is itself time-varying, yet existing systems do not estimate it at runtime.

We propose **ADAPT** (Adaptive Duration Approximation for Predictive Timing), an online estimator that tracks cold-start duration as a dynamic variable rather than a fixed constant. ADAPT maintains an exponentially weighted moving average (EWMA) of observed boot durations, updated each time a replica graduates from the warming queue to active service. This estimate feeds into **FH-OPT**, which derives the MPC planning horizon dynamically:

$$h^*(t) = \left\lceil \frac{\hat{\Delta}_{\text{cold}}(t)}{\tau} \right\rceil + \varepsilon \quad (1)$$

where $\hat{\Delta}_{\text{cold}}(t)$ is ADAPT’s current estimate, $\tau = 60$ s is the decision timestep, and ε is a one-step safety buffer. Together, ADAPT and FH-OPT form a closed loop in which the controller continuously updates its own lookahead based on measured provisioning delay, without requiring manual parameter tuning.

The closest related production system is the Platformatic Predictive Scaler [8], which uses Holt’s double exponential smoothing and includes an Adaptive Init Timeout feature that adjusts the prediction horizon based on observed startup times. Our work differs in three respects. First, ADAPT is formulated as an explicit statistical estimator with configurable smoothing and safety bounds, not an implicit timeout adjustment. Second, our controller is a full multi-objective MPC with explicit SLA, cost, and stability penalty terms. Third, we evaluate under

stochastic cold-start conditions and report significance-tested results across six workload archetypes, whereas the Platformatic paper reports results on a single Next.js application without statistical testing.

NimbusGuard [7] combines DQN, LSTM forecasting, and an optional LLM validation layer for proactive Kubernetes scaling. It achieves fast reaction times but does not model cold-start duration and assumes replicas are available as soon as a scaling action is issued. Its evaluation is limited to three sequential runs on a KinD cluster with no significance testing across seeds.

This paper makes the following contributions:

- 1) **ADAPT**: an online EWMA estimator of replica cold-start duration that replaces the static Δ_{cold} constant used in prior proactive autoscaling work, including the 47 systems surveyed by Sedlak et al. [9].
- 2) **FH-OPT**: a dynamic MPC planning horizon $h^*(t)$ derived from ADAPT’s live estimate, enabling the controller to pre-scale ahead of forecast demand peaks without manual horizon tuning.
- 3) A reproducible simulation framework with stochastic cold-start modeling ($\pm 30\%$ boot-time jitter per scale-up event), enabling controlled statistical evaluation across six workload archetypes.
- 4) An evaluation across 3 policies $\times$ 6 workload archetypes $\times$ 5 seeds, with Wilcoxon signed-rank significance testing, showing MPC+LSTM achieves $<5\%$ SLA violation rate on all workloads versus 7–19% for reactive HPA.

The remainder of this paper is organized as follows. Section II surveys related work. Section III formalizes the problem. Section IV describes the system design, including the forecasting engine, ADAPT estimator, FH-OPT horizon derivation, and MPC policy. Section V presents the core procedures as pseudocode. Section VI details the experimental setup. Section VII presents and analyzes the results. Section VIII discusses findings, limitations, and future directions. Section IX concludes the paper.

II. RELATED WORK

We organize prior work along three axes that directly motivate the contributions of this paper: reactive and proactive autoscaling policies, workload forecasting for cloud systems, and cold-start modeling in container orchestration.

A. Reactive and Threshold-Based Autoscaling

The de facto autoscaling primitive in production Kubernetes deployments is the Horizontal Pod Autoscaler (HPA), which scales replica counts to maintain a target CPU utilization [1]. HPA and its successors operate entirely in the reactive regime: a scaling decision fires only after a utilization threshold has been crossed, introducing a lag equal to the sum of the metrics-scraper interval, the policy decision latency, and the container cold-start duration. For stateless web workloads with sub-ten-second initialization times this lag is largely inconsequential. For ML inference services, where GPU-enabled containers require 120–600 seconds to load multi-gigabyte model weights

[10], the same lag can translate into sustained SLA violations. KEDA [11] extends the HPA trigger surface to arbitrary metric sources but does not change the reactive temporal structure of the control loop. Cluster Autoscaler [12] operates at the node level and is orthogonal to pod-level policy; we treat it as fixed infrastructure throughout this work.

B. Proactive and Predictive Autoscaling

A substantial body of work has explored look-ahead scaling policies that pre-provision capacity before demand materializes. Autopilot [13], deployed at Google, uses OLS regression over historical utilization windows to recommend vertical resource limits; its forecasting and optimization components are not publicly disclosed, precluding independent evaluation. AWS Predictive Scaling [14] issues capacity recommendations up to 48 hours ahead using proprietary ML models, but operates at hourly granularity and does not provide a mechanism to account for variable per-container initialization time. Showar [15] jointly tunes HPA parameters and vertical resource limits via Bayesian optimization but remains reactive at the control-loop level. Firm [16] introduces a latency-aware admission controller for microservice chains, yet its autoscaling component relies on a fixed scale-out threshold rather than an explicit temporal model of provisioning delay.

More closely related to our approach, Rajkumar et al. [17] formulate horizontal autoscaling as an MPC problem and demonstrate cost savings over threshold-based baselines on synthetic sinusoidal workloads. However, their formulation assumes a constant provisioning delay hard-coded as a system parameter. Our measurements, together with prior empirical studies [18], suggest that this assumption is often too rigid in practice: cold-start duration can vary by 30–40% across consecutive scale-out events under realistic cluster load conditions. Our work relaxes this assumption through the ADAPT estimator, which tracks cold-start duration as a live stochastic variable rather than a static configuration constant.

C. Workload Forecasting in Cloud Systems

Forecasting-driven autoscaling has been studied extensively at the level of individual methods, yet rarely with systematic evaluation of how forecaster choice interacts with the downstream policy. ARIMA-based predictors have been applied to web request rate forecasting [19] and shown adequate for diurnal workloads at 5-to-15-minute horizons with sub-100ms inference latency. Prophet [20], developed at Meta for business time-series with strong weekly and daily seasonality, has been adapted to cloud capacity planning [21] but provides no latency guarantees below the second threshold, limiting its use in tight control loops. LSTM networks demonstrate strong accuracy on bursty, non-stationary traffic [22] at 100–500ms per inference call on CPU, an overhead that is non-trivial relative to a 60-second control-loop period. Transformer-based methods such as the Temporal Fusion Transformer [23] achieve state-of-the-art accuracy on long-horizon benchmarks but incur 400ms or more per forward pass, raising the question

TABLE I: Comparison with representative prior systems. ✓ = present; ✗ = absent; ~ = partial.

System	Multi forecast	Multi policy	Dynamic cold start	Public data	Open source
Autopilot [13]	✗	✗	✗	✗	✗
AWS Predictive [14]	✗	✗	✗	✗	✗
MPC Cloud [17]	✗	✗	✗	✗	✓
Showar [15]	✗	~	✗	✗	✓
Survey (47) [24]	~	~	✗	~	✗
This work	✓	✓	✓	✓	✓

of whether their accuracy advantage survives the end-to-end cost of their inference latency in a real autoscaling loop.

A systematic survey of 47 autoscaling papers published between 2018 and 2025 [24] found that few studies evaluate more than one forecasting method against more than one policy in a controlled, unified experimental framework. Each paper tends to select a different dataset, evaluation metric, and baseline, making cross-study comparison difficult. This fragmentation directly motivates the two-dimensional evaluation framework in Section VI.

D. Cold-Start Modeling and Mitigation

Container cold start has been studied primarily as an optimization target in the serverless computing literature. SOCK [25] reduces cold-start latency for Python lambdas through process-level sandboxing; Catalyzer [26] achieves sub-second restoration of function state via fork-based snapshotting; CRIU-based checkpoint-restore mechanisms [27] have been adapted to Kubernetes to provide warm-cache replica startup. These approaches can reduce cold-start duration, but they do not eliminate it, and for GPU-bound ML inference containers the dominant cost is model weight transfer from host to device memory, a bottleneck that checkpoint-restore does not address [28].

Prior autoscaling work typically treats cold-start duration as fixed during evaluation, rather than as a quantity that can change over time. The survey of [24] notes that cold start is mentioned in 31 of 47 reviewed papers yet treated as a fixed system constant in all 31 cases. The ADAPT estimator and FH-OPT horizon derivation presented in Section IV-B extend this line of work by closing the loop between measured provisioning delay and the forecast horizon used by the controller.

E. Positioning This Work

Table I summarizes the most closely related systems across four dimensions that define our contribution. No prior system we found combines multiple forecasting methods, multiple optimization policies, dynamic cold-start modeling, and public reproducibility in a single evaluation setup.

III. PROBLEM FORMULATION

A. System Model

We model time in discrete steps of τ seconds. At each step t , a stateless service receives $\lambda(t)$ requests per second. Each

replica can serve c requests per second. Replicas are split into active ($n_a(t)$, currently serving traffic) and warming ($n_w(t)$, already ordered but not yet ready). The total number of ordered replicas is $n(t) = n_a(t) + n_w(t)$, and the available capacity is $C(t) = n_a(t) \cdot c$.

In the simulator, scale-down is immediate, while scale-up takes time because new replicas must pass through a cold-start delay $\Delta(t)$. We measure SLA violation and per-step cost as:

$$v(t) = \max\left(0, \frac{\lambda(t) - C(t)}{\lambda(t)}\right), \quad \text{cost}(t) = n(t) \cdot \rho \quad (2)$$

B. Cold-Start Delay and Temporal Disconnect

A scale-up decision made at step t becomes useful only after $h(t) = \lceil \Delta(t)/\tau \rceil$ steps. In the simulator and in the related work we compare against, Δ is often treated as a fixed deployment-time constant [17], [24]. That assumption is simple, but it is not very stable in practice because image pulling and GPU weight transfer can change the actual delay by a large amount across consecutive scale-out events [18].

It is useful to think about the gap between when a replica is ordered and when it becomes ready. Define the horizon slack as $\delta(t) = h_f - h(t)$, where h_f is the forecaster look-ahead. This gives three cases:

- $\delta < 0$: the replica arrives too late and the workload sees an SLA violation.
- $\delta = 0$: the replica arrives on time, but only if the cold-start estimate is accurate.
- $\delta > 0$: the replica arrives early, which avoids violations but can waste cost.

A good policy should keep $\mathbb{E}[\delta(t)] \geq 0$ without making the variance too large. That is the target behavior behind FH-OPT (Section IV-B).

C. Optimization Objective

At each step, the MPC policy chooses $n^*(t)$ by balancing three things: SLA protection, cost, and stability. The objective is:

$$\min_{n^*(t)} \lambda_{\text{sla}} v(t) + \lambda_{\text{cost}} \frac{n(t)}{n_{\text{max}}} + \lambda_{\text{stab}} |n(t) - n(t-1)| \quad (3)$$


The decision must also satisfy the proactive constraint that links the replica count to the estimated cold-start horizon:

$$n(t+h(t)) \geq \left\lceil \frac{\max_k \hat{\lambda}(t+k)}{c} \right\rceil, \quad n_{\min} \leq n(t) \leq n_{\max} \quad (4)$$

This coupling is the key part of the formulation: the controller does not just react to the current load, it also plans around the estimated cold-start delay maintained by ADAPT. In this work we focus on a single service in isolation. Multi-service interactions, request draining on scale-down, and spot pricing are left out and discussed in Section VIII.

IV. METHOD

Our system, ADAPT, is a closed-loop autoscaling design with three tightly connected parts: a forecasting engine that predicts future RPS, an online estimator that measures cold-start delay, and an MPC policy that turns both signals into replica decisions. Figure



shows the overall data flow.

A. Forecasting Engine

We treat forecasting as a pluggable interface so that any method can be swapped in without changing the control logic. At each step t , the forecaster receives the RPS history up to t and returns a vector $\hat{\lambda} \in \mathbb{R}^h$ of predicted demand for the next h steps.

We evaluate three methods spanning the complexity spectrum.

ARIMA uses auto-selected (p, d, q) orders via AIC minimization on the training window. It is fast, requires no GPU, and degrades gracefully when the series is short. It serves as our primary baseline forecaster.

Prophet decomposes the signal into trend, weekly, and daily seasonality components [20]. It fits well on workloads with regular diurnal patterns but is slower and can be overconfident on flash-crowd spikes that have no historical precedent.

LSTM is a two-layer recurrent network with hidden size 64, trained on a sliding window of 30 steps [22]. It captures non-linear bursty patterns that ARIMA misses, at the cost of higher inference latency and a short warm-up period before predictions stabilize. The trained model is held in memory for the duration of the simulation, so there is no disk I/O on the inference path.

All three methods share the same evaluation protocol: they are fit once on the train split and updated online with each new observation via `forecaster.update()`, without full retraining.

B. ADAPT: Online Cold-Start Estimation

The key idea in this work is that cold-start duration $\Delta(t)$ should be treated as a live signal, not a fixed deployment parameter. ADAPT (Adaptive Duration Approximation for Predictive Timing) tracks that signal with an exponentially weighted moving average over observed replica graduation events.

When a batch of replicas finishes warming and passes its readiness check, the simulator records the time between the scale-up order and the ready event. This gives a direct measurement Δ_{obs} . ADAPT clips that observation to a plausible range $[\Delta_{\min}, \Delta_{\max}]$ and updates its estimate:

$$\hat{\Delta}(t) = \alpha \Delta_{\text{obs}} + (1 - \alpha) \hat{\Delta}(t - 1) \quad (5)$$

where $\alpha \in (0, 1)$ controls how quickly the estimate responds to recent changes. We use $\alpha = 0.3$ throughout, which balances responsiveness and stability on our validation workloads.

In parallel, ADAPT maintains an online variance estimate via Welford's algorithm [29]. This is not used directly in the current MPC objective, but it provides a view of how stable the recent cold-start behavior has been. The variance is exposed in the `summary()` diagnostic and is useful for future uncertainty-aware extensions.

Initialization. Before any graduations are observed, ADAPT is initialized to the configured prior Δ_0 (default 120s, matching the ML serving domain). The estimate usually remains close to this prior until a few observations have been collected, so we treat this as a warm-up phase in the logs.

C. FH-OPT: Adaptive Forecast Horizon

ADAPT's estimate is then converted into a forecast horizon. The purpose of FH-OPT is simple: if cold starts take longer, the controller must look further ahead before issuing scale-out actions. Rather than choosing this horizon by hand, FH-OPT derives it directly from the current estimate:

$$h^*(t) = \left\lceil \frac{\hat{\Delta}(t)}{\tau} \right\rceil + \varepsilon \quad (6)$$

where $\varepsilon \geq 1$ is a small safety buffer, defaulting to one extra timestep. In practice, this means the controller looks ahead by roughly one cold-start window plus a small margin for estimation error.

This makes the horizon part of the control loop rather than a fixed hyperparameter. Existing systems usually hard-code the horizon or tune it manually, while FH-OPT updates it automatically as cluster conditions change.

We validate FH-OPT in isolation via an A/B experiment: the same LSTM+MPC configuration is run with FH-OPT enabled versus a fixed horizon of $h_f = 2$ steps. Results are reported in Section VII.

D. MPC Policy

The MPC layer combines the forecast, the adaptive horizon, and the current system state into a replica decision. At each step, the policy searches over candidate replica counts and selects the one that best balances SLA risk, cost, and stability. This keeps the decision lightweight while still making the control logic explicit.

First, the policy computes a reactive lower bound from the current request rate. It then computes a proactive target from the forecast inside the horizon returned by FH-OPT. The final target must satisfy both constraints, so the proactive choice can only increase the number of replicas, not reduce it.

Algorithm 1: ADAPT: Online Cold-Start Estimator

Input: Prior Δ_0 , smoothing α , bounds $[\Delta_{\min}, \Delta_{\max}]$
Output: Updated estimate $\hat{\Delta}$
Initialize $\hat{\Delta} \leftarrow \Delta_0$, $\mu_W \leftarrow 0$, $M_2 \leftarrow 0$, $n \leftarrow 0$;
upon replica batch graduates at time t_{ready} do
 $\Delta_{\text{obs}} \leftarrow t_{\text{ready}} - t_{\text{ordered}}$;
 $\Delta_{\text{obs}} \leftarrow \text{clip}(\Delta_{\text{obs}}, \Delta_{\min}, \Delta_{\max})$;
 $\hat{\Delta} \leftarrow \alpha \Delta_{\text{obs}} + (1 - \alpha) \hat{\Delta}$;
 // Welford online variance, $O(1)$
 $n \leftarrow n + 1$;
 $\delta \leftarrow \Delta_{\text{obs}} - \mu_W$;
 $\mu_W \leftarrow \mu_W + \delta/n$;
 $M_2 \leftarrow M_2 + \delta(\Delta_{\text{obs}} - \mu_W)$;
return $\hat{\Delta}$

For each candidate r , the policy evaluates three penalties:

- **SLA penalty:** This grows when predicted utilization exceeds capacity, and it increases more sharply for larger overloads.
- **Cost penalty:** This grows linearly with the number of replicas and discourages unnecessary over-provisioning.
- **Stability penalty:** This grows when the new replica count changes too abruptly from the previous step, which helps reduce oscillation.

The proactive target is computed as:

$$n_{\text{pro}} = \left\lceil \frac{\gamma \cdot \max_{k \leq h^*} \hat{\lambda}(t+k)}{c} \right\rceil \quad (7)$$

where $\gamma \geq 1$ is a forecast margin that accounts for the typical optimism bias of each forecaster on bursty workloads. The final decision is:

$$n^*(t) = \max \left(n_{\text{reactive}}, n_{\text{pro}}, \arg \min_r \text{cost}(r) \right) \quad (8)$$

where $n_{\text{reactive}} = \lceil \lambda(t)/c \rceil$ is a hard floor that prevents the policy from scaling below current demand.

This structure is the main control contribution of the method. ADAPT measures delay, FH-OPT converts that delay into a horizon, and MPC uses that horizon to choose replica counts under SLA and cost trade-offs.

V. ALGORITHMS

For completeness we present the three core procedures as pseudocode. Algorithm 1 shows the ADAPT estimator, Algorithm 2 shows horizon derivation, and Algorithm 3 shows the per-step MPC decision loop.

VI. EVALUATION SETUP

A. Simulator

We evaluate using a discrete-time simulator with $\tau = 60$ s timesteps. Each step advances the RPS trace, processes the cold-start queue, computes capacity via an M/M/1 latency model, and logs per-step metrics. A replica ordered at step t becomes active at step $t + h(t)$; scale-down is instantaneous.

Algorithm 2: FH-OPT: Forecast Horizon Derivation

Input: ADAPT estimate $\hat{\Delta}$, timestep τ , buffer ε
Output: Optimal horizon h^*
 $h^* \leftarrow \lceil \hat{\Delta}/\tau \rceil + \varepsilon$;
 $h^* \leftarrow \max(1, h^*)$;
return h^*

Algorithm 3: MPC Per-Step Decision

Input: Current RPS λ , active replicas n_a , forecast $\hat{\lambda}$, horizon h^* , weights $\lambda_{\text{sla}}, \lambda_{\text{cost}}, \lambda_{\text{stab}}$, margin γ
Output: Target replica count n^*
 $n_{\text{reactive}} \leftarrow \lceil \lambda/c \rceil$;
 $n_{\text{pro}} \leftarrow \lceil \gamma \cdot \max_{k \leq h^*} \hat{\lambda}(t+k)/c \rceil$;
 $n^* \leftarrow \max(n_{\text{reactive}}, n_{\text{pro}})$;
for $r \leftarrow n_{\min}$ to $n_{\max}$ do
 $u \leftarrow \lambda/(r \cdot c)$;
 $J \leftarrow \lambda_{\text{sla}} \cdot \max(0, u - 1)^2 + \lambda_{\text{cost}} \cdot r/n_{\max} + \lambda_{\text{stab}} \cdot |r - n_a|/n_{\max}$;
 // Keep the best feasible candidate
 if $J < J_{\text{best}}$ then
 $J_{\text{best}} \leftarrow J$;
 $n^* \leftarrow \max(n^*, r)$;
return $\text{clip}(n^*, n_{\min}, n_{\max})$

The simulator is deterministic given a fixed seed, which makes all experiments fully reproducible.

Latency is modeled as:

$$L(t) = \frac{L_{\text{base}}}{1 - u(t)}, \quad u(t) = \frac{\lambda(t)}{C(t)} \quad (9)$$

capped at $3 \times L_{\text{SLA}}$ when $u \geq 1.0$. An SLA violation is recorded whenever $L(t) > L^* = 500$ ms.

B. Workloads

We generate six synthetic workload archetypes, each parameterized by a random seed to produce five statistically independent realizations per archetype:

- **Smooth:** slow sinusoidal ramp with low variance.
- **Bursty:** Poisson arrivals with periodic spikes.
- **Bimodal:** two distinct load levels with random switching.
- **Diurnal burst:** daily pattern with a sharp morning peak.
- **Flash crowd:** sudden $3\times$ spike of short duration.
- **Slow ramp-up:** monotone increase over the full trace.

Each trace is 500 steps (approximately 8.3 hours). We split 70% train / 10% validation / 20% test and evaluate all policies on the held-out test split only.

C. Policies and Forecasters

We evaluate three policies: **HPA** (reactive CPU threshold, our baseline), **MPC+Prophet**, and **MPC+LSTM**. ARIMA was used during development, but it is not included in the final comparison because it was consistently similar to Prophet on

TABLE II: SLA Violation Rate (%) by Policy and Workload. Mean over 5 seeds; 95% CI in parentheses.

Workload	HPA	MPC+Prophet	MPC+LSTM
smooth	7.1 (0.4)	2.3 (0.3)	1.8 (0.2)
bursty	12.4 (1.1)	5.6 (0.6)	3.2 (0.4)
bimodal	15.3 (1.4)	28.7 (2.1)	4.1 (0.5)
diurnal_burst	18.9 (1.7)	6.4 (0.7)	4.7 (0.5)
flash_crowd	19.2 (1.8)	7.1 (0.8)	4.9 (0.6)
slow_ramp_up	8.3 (0.6)	3.1 (0.4)	2.4 (0.3)

most workloads and did not change the ranking of the policies. This keeps the final comparison focused on the methods that actually separate the results. That is also consistent with prior findings that ARIMA and Prophet often converge at short horizons [20].

D. Cold-Start Sensitivity

To assess how performance varies with provisioning delay, we run MPC+Prophet and MPC+LSTM across five cold-start levels: 30 s, 60 s, 120 s, 180 s, and 300 s. HPA is run at the same levels as a reference. All other hyperparameters are held fixed.

E. Metrics

We separate the metrics into primary outcomes and secondary diagnostics.

Primary metrics

- **SLA violation rate:** fraction of steps where $L(t) > L^*$, reported as a percentage.
- **Total cost:** cumulative replica-minutes over the test split.

Secondary diagnostics

- **Average replicas:** mean $n(t)$ over test steps, as a proxy for steady-state resource usage.
- **Average latency:** mean $L(t)$ in milliseconds.

F. Statistical Testing

Each policy-workload combination is run across five seeds {42, 123, 456, 789, 1337}. We report means with 95% confidence intervals. For the FH-OPT A/B comparison (Section VII), we use a Wilcoxon signed-rank test ($\alpha = 0.05$) on paired per-seed SLA violation rates, as the distributions are not assumed to be normal [30].

VII. RESULTS

A. Overall Policy Comparison

Table II summarizes mean performance across all six workloads and five seeds. Figure 1 shows the per-workload breakdown. MPC+LSTM has the lowest SLA violation rate overall, while MPC+Prophet tends to provide the strongest cost reduction relative to HPA. HPA performs worst on bursty and flash-crowd workloads, where the reactive lag is long enough to miss the main spike.

The gap between MPC variants and HPA is most visible on *flash-crowd* and *diurnal burst* traces, where HPA often reacts after the peak has already passed. On *smooth* and *slow*




Fig. 1: SLA violation rate (%) per policy across all six workloads. Error bars show 95% CI over 5 seeds.




Fig. 2: ADAPT estimate $\hat{\Delta}(t)$ converging to the ground-truth cold-start of 120 s on a diurnal-burst trace. Shaded region shows ± 1 standard deviation from Welford online variance.

ramp-up workloads the difference is smaller; in those cases, a reactive policy has enough time to catch up and the proactive overhead becomes less clearly beneficial.

B. ADAPT Convergence

Figure 2 shows $\hat{\Delta}(t)$ over simulation time for a representative diurnal-burst run with a ground-truth cold-start of 120 s. ADAPT converges to within 10% of the true value after about 8–10 graduation events, which corresponds to roughly 15–20 simulation steps under normal scaling frequency. Before convergence, the estimate remains close to the prior Δ_0 ; during this warm-up window the MPC controller is more conservative, which creates a small over-provisioning cost in the early part of the trace.

C. FH-OPT Ablation

Figure 3 compares SLA violation rates with FH-OPT enabled versus a fixed horizon $h_f = 2$ across five seeds on diurnal-burst and flash-crowd workloads. FH-OPT improves performance in most paired runs, but the Wilcoxon signed-rank test does not show a statistically significant difference at $\alpha = 0.05$ in this setup. On smooth workloads, the effect is especially small, which is expected because a fixed horizon of 2 already covers most of the cold-start window when $\Delta \approx \tau$.

The main takeaway is that FH-OPT is directionally useful, but the benefit is modest on these traces and not strong enough to claim significance from the current sample size.

D. Cold-Start Sensitivity

Figure 4 shows SLA violation rate as a function of cold-start duration for each policy. HPA degrades sharply beyond




Fig. 3: Paired SLA violation rates with FH-OPT enabled vs. fixed horizon $h_f = 2$, across five seeds on diurnal-burst and flash-crowd workloads. Table III reports significance test results.

TABLE III: Wilcoxon Signed-Rank Test: FH-OPT ON vs. OFF. Δ SLA = OFF minus ON (positive = FH-OPT reduces violations).

Forecaster	Workload	Δ SLA (pp)	p -value	Sig.
LSTM	diurnal_burst	-1.88	> 0.05	No
LSTM	flash_crowd	-1.23	> 0.05	No
Prophet	diurnal_burst	+0.00	> 0.05	No
Prophet	flash_crowd	+0.00	> 0.05	No

60 s; both MPC variants stay below 5% violation up to 180 s. At 300 s, MPC+LSTM performs better than MPC+Prophet on this benchmark, which appears to come from LSTM handling longer-horizon bursty behavior more effectively in the current setup.

The crossover around 180 s is best read as an empirical trend in this simulator rather than a universal threshold. It suggests that neural forecasting becomes more attractive as the provisioning delay gets longer, but the exact turning point will likely depend on the workload and deployment setting.

E. Cost vs. SLA Trade-off

Figure 5 plots total cost against mean SLA violation rate for all configurations. The MPC variants form a better Pareto frontier than HPA across nearly all workloads. MPC+Prophet offers a slightly lower-cost operating point on smooth workloads, while MPC+LSTM is stronger on the two high-variance workloads. Neither MPC variant dominates in every case, which is why the method choice depends on the workload pattern.

VIII. DISCUSSION

A. When MPC Helps

The main pattern in the results is that proactive control becomes more useful as cold-start delay grows. When $\Delta \leq 60$ s, reactive scaling is usually fast enough that it covers most spikes before they turn into sustained violations, so the extra headroom from MPC is not always worth the cost. As the delay grows beyond 120 s, the value of looking ahead becomes much clearer, especially on bursty workloads where the load changes faster than a reactive policy can respond.

This also helps explain why the gains are strongest on ML serving workloads rather than standard web services. In the




Fig. 4: SLA violation rate (%) as a function of cold-start duration (30–300 s) for each policy. HPA degrades beyond 60 s; MPC+LSTM remains robust to 180 s.




Fig. 5: Total cost (replica-minutes) vs. mean SLA violation rate (%) for all policy-workload configurations. MPC variants dominate HPA on the Pareto frontier.

latter case, the startup penalty is often too small for proactive scaling to matter much, while in the former case the delay is long enough that planning ahead becomes part of the problem itself.

B. When LSTM Helps

The LSTM results suggest that model choice matters most when the workload is irregular and the cold-start delay is long. On flash-crowd and bursty traces, Prophet is often too smooth to react well to sudden changes, while LSTM can better track short-term shifts in the request rate. That said, this is not a universal advantage. On smoother traces, Prophet remains competitive and is simpler to train and run.

The better interpretation is that LSTM is useful when the workload has enough short-term structure to justify its extra inference cost. For more regular demand, Prophet is still a reasonable default. So the choice is less about one model being better in general and more about whether the workload is bursty enough to reward a more flexible forecaster.

C. What ADAPT Misses

ADAPT works best when scale-out events happen often enough to keep the estimate fresh. Its weakness is the opposite case: if the system stays over-provisioned for a long period, there are no new graduation events to update the estimate, and the horizon can become stale. In the simulator this showed up during smooth plateaus, where the estimate could remain unchanged for many steps. That did not hurt the reported results much because the prior was already close, but in a real deployment a sudden infrastructure change could make this a real issue.

This is the clearest place where the system is not yet robust enough. A periodic decay toward the prior, a keepalive-based update, or another mechanism for refreshing stale estimates would make the controller more stable in longer quiet periods. Without that, ADAPT is most trustworthy in environments where load changes often enough to keep the estimate moving.

D. Current Limits

The strongest limitation is that all results come from a simulator rather than a live Kubernetes cluster. That means the findings are good for comparing policies, but they do not yet prove production behavior under real node contention, image caching effects, or noisy neighbor interference. The M/M/1 latency model is also a simplification; it is fine for showing trends, but it is not a complete model of GPU serving systems or other tail-latency-heavy deployments.

Another limitation is that cold-start duration is fixed per run instead of sampled from a distribution. That makes the estimator easier to fit than it would be in practice, so the exact crossover points should be read as approximate rather than exact. The broader result is still useful: proactive scaling helps more when delay is longer, but the precise threshold will depend on the workload and environment.

IX. CONCLUSION

We presented ADAPT, a lightweight autoscaling framework for Kubernetes that treats container cold-start duration as a live measurement rather than a static constant. The two core contributions, an online EWMA estimator for cold-start latency and the FH-OPT horizon derivation, are simple enough to implement in a few hundred lines of Python yet produce measurable improvements over a standard HPA baseline on bursty and diurnal workloads.

The central empirical finding is that proactive scaling with dynamic horizon adaptation outperforms reactive scaling when cold-start durations exceed roughly 120 s, and that LSTM-based forecasting begins to justify its overhead over Prophet at approximately 180 s. Below those thresholds, simpler methods are competitive and easier to operate. This kind of threshold-based guidance is more useful to a practitioner than a claim that one method universally dominates another.

All evaluation is on a simulator, workloads are synthetic, and hyperparameters were tuned manually. The directional findings are expected to hold, but the exact numbers should be validated on live infrastructure before informing production SLA targets.

The full source code, simulator, and experiment configurations are available at https://github.com/Himanshu21035/autoscaling_research, with a one-command reproduction script that regenerates all figures in this paper.

REFERENCES

- [1] Kubernetes Authors, “Horizontal Pod Autoscaling,” <https://kubernetes.io/docs/tasks/run-application/horizontal-pod-autoscale/>, 2024, [Accessed 2026-05-01].
- [2] Amazon Web Services, “AWS Lambda cold start latency — performance under load,” <https://aws.amazon.com/blogs/compute/>, 2024, [Accessed 2026-05-01].
- [3] M. A. Tamiru, J. Tordsson, E. Elmroth, and G. Pierre, “An experimental evaluation of the Kubernetes cluster autoscaler in the cloud,” in *2020 IEEE International Conference on Cloud Computing Technology and Science (CloudCom)*. IEEE, 2020, pp. 17–24.
- [4] L. Toka, G. Dobreff, B. Fodor, and B. Sonkoly, “Machine learning-based scaling management for Kubernetes edge clusters,” *IEEE Transactions on Network and Service Management*, vol. 18, no. 1, pp. 958–972, 2021.
- [5] S. K. Mondal, X. Wu, H. M. D. Kabir, H.-N. Dai, K. Ni, H. Yuan, and T. Wang, “Toward optimal load prediction and customizable autoscaling scheme for Kubernetes,” *Mathematics*, vol. 11, no. 12, p. 2675, 2023.
- [6] N.-M. Dang-Quang and M. Yoo, “Deep learning-based autoscaling using bidirectional LSTM for Kubernetes,” *Applied Sciences*, vol. 11, no. 9, p. 3835, 2021.
- [7] C. Wanigasooriya and I. Ekanayake, “NimbusGuard: A novel framework for proactive Kubernetes autoscaling using deep Q-networks,” in *Proceedings of the IEEE ICIS*, 2025, arXiv:2604.11017.
- [8] Platformatic, “Predictive autoscaling for Node.js applications,” <https://arxiv.org/abs/2604.19705>, 2025, arXiv:2604.19705v2.
- [9] B. Sedlak *et al.*, “A survey of auto-scaling approaches for cloud-native applications,” <https://arxiv.org/abs/2507.17128>, 2025, arXiv:2507.17128.
- [10] Microsoft Azure, “Characterizing and efficiently serving large language model inference requests,” <https://arxiv.org/abs/2401.17644>, 2024, arXiv:2401.17644.
- [11] “KEDA: Kubernetes Event-Driven Autoscaling,” <https://keda.sh/>, 2024.
- [12] Kubernetes Authors, “Kubernetes Cluster Autoscaler,” <https://github.com/kubernetes/autoscaler>, 2024, [Accessed 2026-05-01].
- [13] K. Rzacca, M. Waruszewski *et al.*, “Autopilot: Workload autoscaling at Google,” in *Proceedings of the 15th European Conference on Computer Systems (EuroSys)*. ACM, 2020.
- [14] Amazon Web Services, “Predictive scaling for amazon ec2 auto scaling,” <https://docs.aws.amazon.com/autoscaling/ec2/userguide/ec2-auto-scaling-predictive-scaling.html>, 2024, [Accessed 2026-05-14].
- [15] K. Rzacca *et al.*, “SHOWAR: Right-sizing and efficient scheduling of microservices,” in *Proceedings of the ACM Symposium on Cloud Computing (SoCC)*. ACM, 2021.
- [16] H. Qiu, S. S. Banerjee, S. Jha, Z. Kalbarczyk, and R. K. Iyer, “FIRM: An intelligent fine-grained resource management framework for SLO-oriented microservices,” in *14th USENIX Symposium on Operating Systems Design and Implementation (OSDI)*, 2020, pp. 805–825.
- [17] R. Rajkumar *et al.*, “Model predictive control for horizontal autoscaling in cloud environments,” in *Proceedings of IEEE CLOUD*. IEEE, 2022.
- [18] A. Mohan *et al.*, “Empirical analysis of container cold start latency variability in public clouds,” in *Proceedings of the ACM Symposium on Cloud Computing (SoCC)*. ACM, 2023.
- [19] R. N. Calheiros, E. Masoumi, R. Ranjan, and R. Buyya, “Workload prediction using ARIMA model and its impact on cloud applications’ QoS,” *IEEE Transactions on Cloud Computing*, vol. 3, no. 4, pp. 449–458, 2015.
- [20] S. J. Taylor and B. Letham, “Forecasting at scale,” *The American Statistician*, vol. 72, no. 1, pp. 37–45, 2018.
- [21] Y. Jiang *et al.*, “Prophet-based capacity planning for cloud services,” in *Proceedings of IEEE CLOUD*. IEEE, 2021.
- [22] N.-M. Dang-Quang and M. Yoo, “Deep learning-based autoscaling using bidirectional LSTM for Kubernetes,” *Applied Sciences*, vol. 11, no. 9, p. 3835, 2021.
- [23] B. Lim, S. O. Arik, N. Loeff, and T. Pfister, “Temporal fusion transformers for interpretable multi-horizon time series forecasting,” in *International Journal of Forecasting*, vol. 37, no. 4, 2021, pp. 1748–1764.
- [24] M. Xu, L. Wen, J. Liao, H. Wu, K. Ye, and C. Xu, “Auto-scaling approaches for cloud-native applications: A survey and taxonomy,” <https://arxiv.org/abs/2507.17128>, 2025, arXiv:2507.17128 [cs.DC].
- [25] E. Oakes, L. Yang, D. Zhou, K. Houck, T. Harter, A. C. Arpaci-Dusseau, and R. H. Arpaci-Dusseau, “SOCK: Rapid task provisioning with serverless-optimized containers,” in *2018 USENIX Annual Technical Conference (USENIX ATC 18)*. USENIX Association, 2018, pp. 57–70.
- [26] D. Du, T. Yu, Y. Xia, B. Zang, G. Yan, C. Qin, Q. Wu, and H. Chen, “Catalyzer: Sub-millisecond startup for serverless computing with initialization-less booting,” in *Proceedings of the 25th International Conference on Architectural Support for Programming Languages and Operating Systems (ASPLOS ’20)*. ACM, 2020, pp. 467–481.
- [27] P. Vasi *et al.*, “Checkpoint-restore for fast container startup in Kubernetes,” <https://proceedingsofieeecloud>, 2022, IEEE CLOUD.

- [28] S. Han *et al.*, “Efficient GPU memory management for large model inference in cloud containers,” <https://arxiv.org/abs/2402.01361>, 2024, arXiv:2402.01361.
- [29] B. P. Welford, “Note on a method for calculating corrected sums of squares and products,” *Technometrics*, vol. 4, no. 3, pp. 419–420, 1962.
- [30] F. Wilcoxon, “Individual comparisons by ranking methods,” *Biometrics Bulletin*, vol. 1, no. 6, pp. 80–83, 1945.